

Breast Cancer Detection



GROUP NO. 3

1. Minita Joshee [23]
2. Amaan Nizam [46]
3. Abhiram Pillai [52]
4. Anne Pinto [53]

PROJECT GUIDE : DR. SATISHKUMAR CHAVAN

**DEPARTMENT OF ELECTRONICS AND TELECOMMUNICATION
DON BOSCO INSTITUTE OF TECHNOLOGY**

Objectives

- To classify the tumors as Benign and Malignant tumors.
- To detect location of the tumour.
- To automate the process of detection of abnormal tissues.
- To evaluate the performance of various Deep Learning approaches for detection and segmentation of tumours



Outcomes

- Software tool which automatically detects cancerous tissues.
- Software will be able to identify location of tumour and detect volume of tumour.
- It will act as an assisting tool to radiologists to classify or choose the abnormal mammogram and prioritise based on level of concern.
- Learning how to write a Technical paper using LaTeX

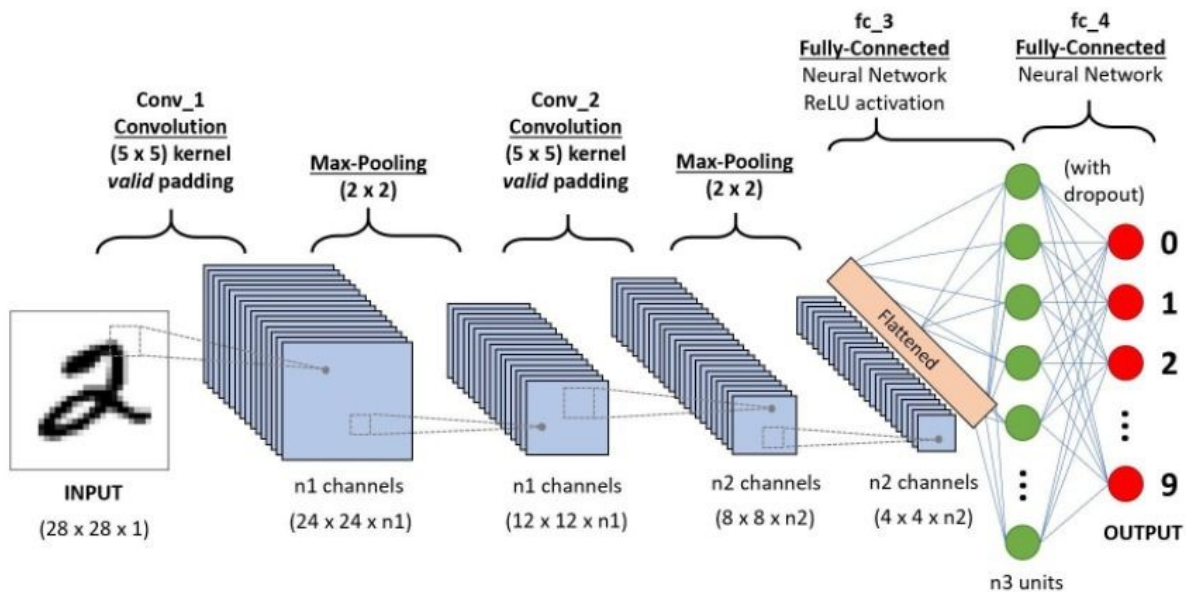


Problem Statement

To develop an automated detection and segmentation of tumours using mammogram in Cranial-Caudal and Medial-lateral oblique (CC and MLO) views using Deep Learning Techniques



CNN




Training Using Alexnet

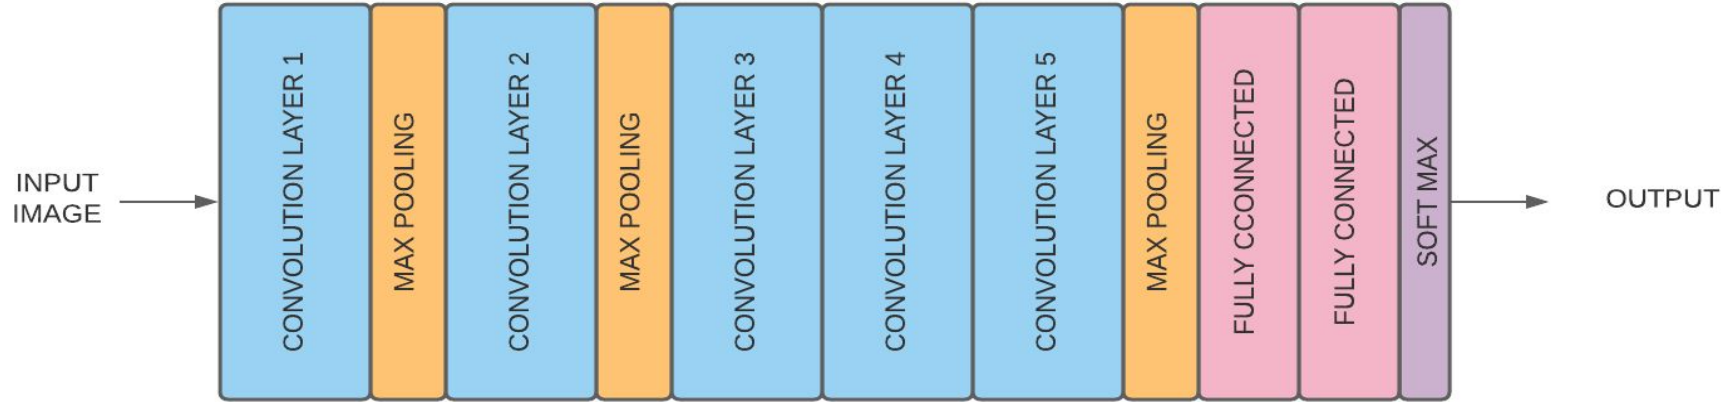
- 8 layers deep neural network
- A total of 60 million parameters
- 5 Convolutional layers and 3 Fully Connected (dense) layers
- Dataset consisting of 15 million images, over 22,000 categories with variable resolution images so scaled down to fixed resolution 256x256
- Why ReLu Non-linearity?



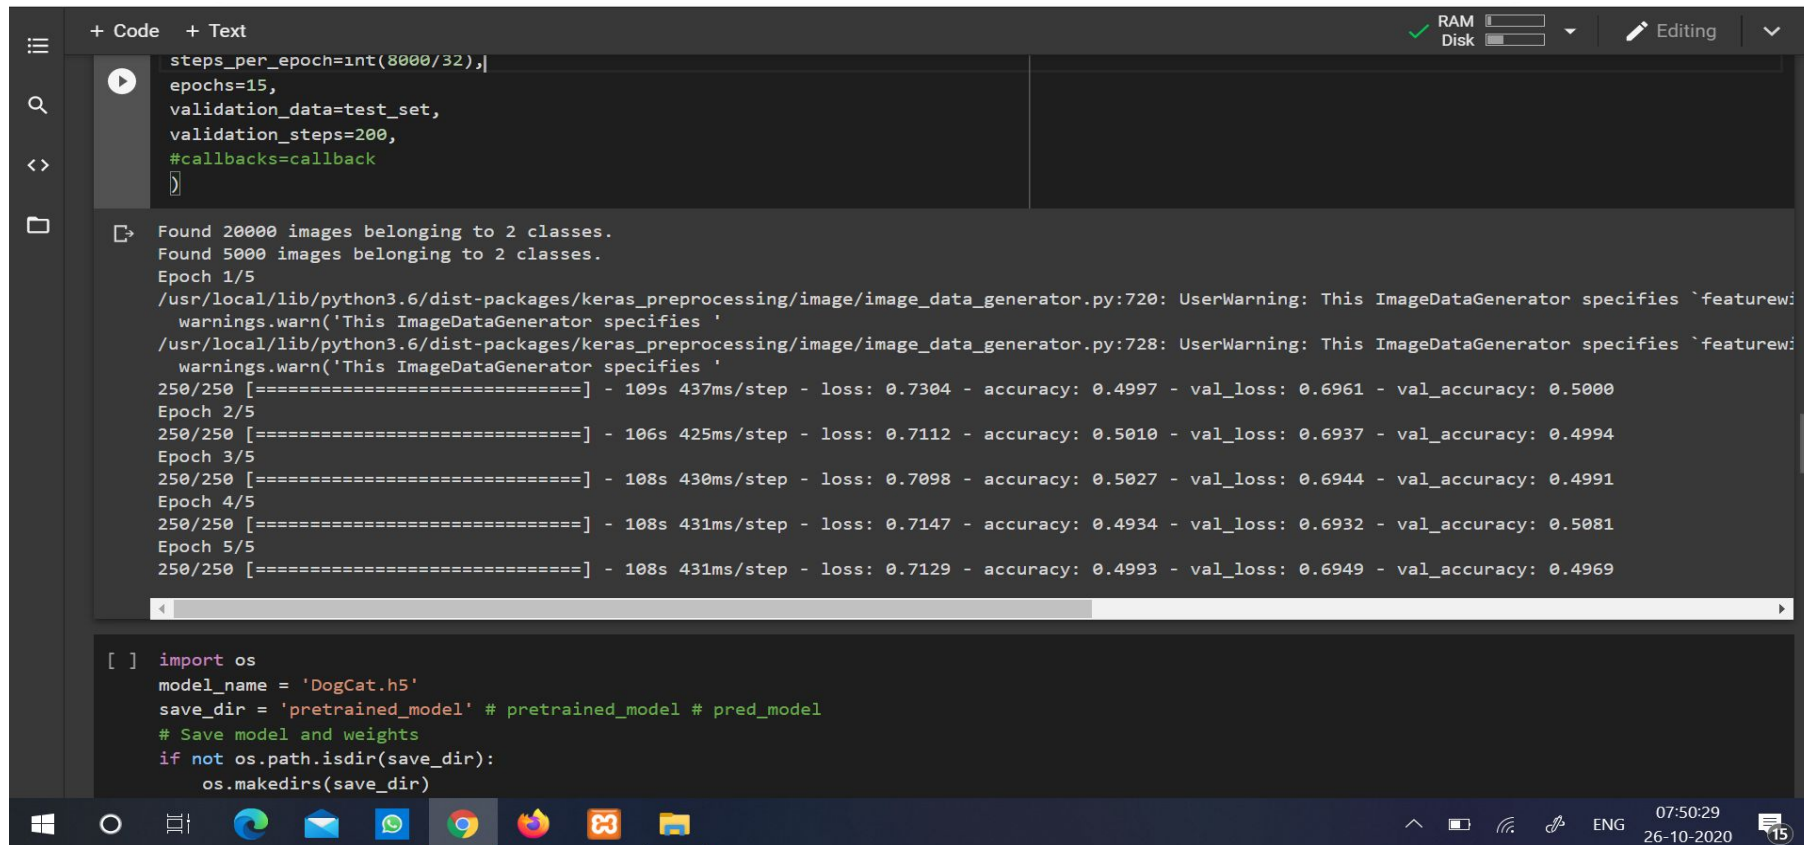
Architecture of Alexnet



Architecture



Results of Alexnet



The screenshot shows a Jupyter Notebook interface with a dark theme. The top bar includes a '+ Code' button, a '+ Text' button, and system status indicators for RAM and Disk. The left sidebar contains icons for a table of contents, search, and file explorer. The main area is divided into two sections: a code editor and an output area.

Code Editor:

```
steps_per_epoch=int(8000/32),|
epochs=15,
validation_data=test_set,
validation_steps=200,
#callbacks=callback
|
```

Output Area:

```
Found 20000 images belonging to 2 classes.
Found 5000 images belonging to 2 classes.
Epoch 1/5
/usr/local/lib/python3.6/dist-packages/keras_preprocessing/image/image_data_generator.py:720: UserWarning: This ImageDataGenerator specifies `featurew
warnings.warn('This ImageDataGenerator specifies '
/usr/local/lib/python3.6/dist-packages/keras_preprocessing/image/image_data_generator.py:728: UserWarning: This ImageDataGenerator specifies `featurew
warnings.warn('This ImageDataGenerator specifies '
250/250 [=====] - 109s 437ms/step - loss: 0.7304 - accuracy: 0.4997 - val_loss: 0.6961 - val_accuracy: 0.5000
Epoch 2/5
250/250 [=====] - 106s 425ms/step - loss: 0.7112 - accuracy: 0.5010 - val_loss: 0.6937 - val_accuracy: 0.4994
Epoch 3/5
250/250 [=====] - 108s 430ms/step - loss: 0.7098 - accuracy: 0.5027 - val_loss: 0.6944 - val_accuracy: 0.4991
Epoch 4/5
250/250 [=====] - 108s 431ms/step - loss: 0.7147 - accuracy: 0.4934 - val_loss: 0.6932 - val_accuracy: 0.5081
Epoch 5/5
250/250 [=====] - 108s 431ms/step - loss: 0.7129 - accuracy: 0.4993 - val_loss: 0.6949 - val_accuracy: 0.4969
```

Code Editor (Bottom):

```
[ ] import os
model_name = 'DogCat.h5'
save_dir = 'pretrained_model' # pretrained_model # pred_model
# Save model and weights
if not os.path.isdir(save_dir):
    os.makedirs(save_dir)
```

The bottom of the image shows a Windows taskbar with various application icons and a system tray displaying the time (07:50:29) and date (26-10-2020).

Training using VGG-16

- VGGNet consists of 16 convolutional layers with only 3x3 kernels.
- **Input:** VGG takes in a 224x224 pixel RGB image.
- **Convolutional Layers:** The convolutional layers in VGG use a very small receptive field (3x3, the smallest possible size that still captures left/right and up/down).
- **Fully-Connected Layers:** VGG has three fully-connected layers: 4096 channels each and the last one has 1000 channels
- **Hidden Layers:** All of VGG's hidden layers use ReLU (a huge innovation from AlexNet that cut training time).



Architecture



Results of VGG-16

```
Epoch 1/10
100/100 [=====] - 33s 333ms/step - loss: 0.6756 - acc: 0.6855 - val_loss: 0.3504 - val_acc: 0.8335
Epoch 2/10
100/100 [=====] - 33s 332ms/step - loss: 0.4837 - acc: 0.7655 - val_loss: 0.4827 - val_acc: 0.7800
Epoch 3/10
100/100 [=====] - 33s 328ms/step - loss: 0.4147 - acc: 0.8110 - val_loss: 0.2349 - val_acc: 0.8980
Epoch 4/10
100/100 [=====] - 33s 330ms/step - loss: 0.4037 - acc: 0.8160 - val_loss: 0.2329 - val_acc: 0.9030
Epoch 5/10
100/100 [=====] - 33s 327ms/step - loss: 0.4103 - acc: 0.8205 - val_loss: 0.2368 - val_acc: 0.8960
Epoch 6/10
100/100 [=====] - 33s 326ms/step - loss: 0.3732 - acc: 0.8255 - val_loss: 0.2103 - val_acc: 0.9070
Epoch 7/10
100/100 [=====] - 33s 332ms/step - loss: 0.3564 - acc: 0.8420 - val_loss: 0.2249 - val_acc: 0.9030
Epoch 8/10
100/100 [=====] - 33s 329ms/step - loss: 0.3532 - acc: 0.8455 - val_loss: 0.2329 - val_acc: 0.9050
Epoch 9/10
100/100 [=====] - 33s 329ms/step - loss: 0.3583 - acc: 0.8440 - val_loss: 0.2276 - val_acc: 0.9015
Epoch 10/10
100/100 [=====] - 33s 331ms/step - loss: 0.3418 - acc: 0.8475 - val_loss: 0.2320 - val_acc: 0.9045
```



Training using GoogleNet

- 22 layers deep.
- The architecture was designed to keep computational efficiency in mind. The idea behind that the architecture can be run on individual devices even with low computational resources.
- Uses many different methods such as 1×1 convolution and global average pooling that enables it to create a deeper architecture.

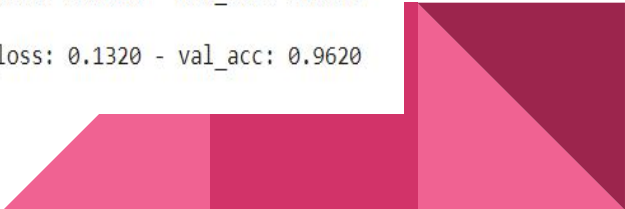


Architecture

| type | patch size/ stride | output size | depth | #1×1 | #3×3 reduce | #3×3 | #5×5 reduce | #5×5 | pool proj | params | ops |
|----------------|-----------------------|----------------|-------|------|----------------|------|----------------|------|--------------|--------|------|
| convolution | 7×7/2 | 112×112×64 | 1 | | | | | | | 2.7K | 34M |
| max pool | 3×3/2 | 56×56×64 | 0 | | | | | | | | |
| convolution | 3×3/1 | 56×56×192 | 2 | | 64 | 192 | | | | 112K | 360M |
| max pool | 3×3/2 | 28×28×192 | 0 | | | | | | | | |
| inception (3a) | | 28×28×256 | 2 | 64 | 96 | 128 | 16 | 32 | 32 | 159K | 128M |
| inception (3b) | | 28×28×480 | 2 | 128 | 128 | 192 | 32 | 96 | 64 | 380K | 304M |
| max pool | 3×3/2 | 14×14×480 | 0 | | | | | | | | |
| inception (4a) | | 14×14×512 | 2 | 192 | 96 | 208 | 16 | 48 | 64 | 364K | 73M |
| inception (4b) | | 14×14×512 | 2 | 160 | 112 | 224 | 24 | 64 | 64 | 437K | 88M |
| inception (4c) | | 14×14×512 | 2 | 128 | 128 | 256 | 24 | 64 | 64 | 463K | 100M |
| inception (4d) | | 14×14×528 | 2 | 112 | 144 | 288 | 32 | 64 | 64 | 580K | 119M |
| inception (4e) | | 14×14×832 | 2 | 256 | 160 | 320 | 32 | 128 | 128 | 840K | 170M |
| max pool | 3×3/2 | 7×7×832 | 0 | | | | | | | | |
| inception (5a) | | 7×7×832 | 2 | 256 | 160 | 320 | 32 | 128 | 128 | 1072K | 54M |
| inception (5b) | | 7×7×1024 | 2 | 384 | 192 | 384 | 48 | 128 | 128 | 1388K | 71M |
| avg pool | 7×7/1 | 1×1×1024 | 0 | | | | | | | | |
| dropout (40%) | | 1×1×1024 | 0 | | | | | | | | |
| linear | | 1×1×1000 | 1 | | | | | | | 1000K | 1M |
| softmax | | 1×1×1000 | 0 | | | | | | | | |

Results of GoogleNet

```
Epoch 1/10
100/100 [=====] - 28s 278ms/step - loss: 1.1381 - acc: 0.8415 - val_loss: 0.1191 - val_acc: 0.9595
Epoch 2/10
100/100 [=====] - 26s 256ms/step - loss: 0.4630 - acc: 0.8785 - val_loss: 0.3551 - val_acc: 0.9185
Epoch 3/10
100/100 [=====] - 27s 267ms/step - loss: 0.3491 - acc: 0.9020 - val_loss: 0.1194 - val_acc: 0.9565
Epoch 4/10
100/100 [=====] - 26s 256ms/step - loss: 0.3000 - acc: 0.9140 - val_loss: 0.1213 - val_acc: 0.9640
Epoch 5/10
100/100 [=====] - 26s 265ms/step - loss: 0.3183 - acc: 0.9070 - val_loss: 0.1598 - val_acc: 0.9550
Epoch 6/10
100/100 [=====] - 26s 256ms/step - loss: 0.3174 - acc: 0.9205 - val_loss: 0.1801 - val_acc: 0.9575
Epoch 7/10
100/100 [=====] - 27s 266ms/step - loss: 0.3334 - acc: 0.9105 - val_loss: 0.1175 - val_acc: 0.9580
Epoch 8/10
100/100 [=====] - 26s 257ms/step - loss: 0.3147 - acc: 0.9080 - val_loss: 0.2037 - val_acc: 0.9440
Epoch 9/10
100/100 [=====] - 26s 263ms/step - loss: 0.2794 - acc: 0.9185 - val_loss: 0.2248 - val_acc: 0.9445
Epoch 10/10
100/100 [=====] - 26s 256ms/step - loss: 0.3136 - acc: 0.9145 - val_loss: 0.1320 - val_acc: 0.9620
```




EfficientNet

- Alexnet \longrightarrow VGG-16 \longrightarrow Googlenet.
- Process of scaling up CNNs is not well understood; many ways of doing this, mostly arbitrarily chosen.
- Three types of scaling.- width,depth,resolution.
- Efficientnet uses compound scaling.
- Saturation happens if only one type of scaling is done.



EfficientNet



EfficientNet

```
Epoch 1/10
100/100 [=====] - 32s 324ms/step - loss: 0.3151 - accuracy: 0.9055 - val_loss: 0.0549 - val_accuracy: 0.9670
Epoch 2/10
100/100 [=====] - 30s 304ms/step - loss: 0.3398 - accuracy: 0.9160 - val_loss: 0.0211 - val_accuracy: 0.9670
Epoch 3/10
100/100 [=====] - 30s 303ms/step - loss: 0.4132 - accuracy: 0.9070 - val_loss: 0.0252 - val_accuracy: 0.9780
Epoch 4/10
100/100 [=====] - 30s 303ms/step - loss: 0.3304 - accuracy: 0.9170 - val_loss: 0.5167 - val_accuracy: 0.9810
Epoch 5/10
100/100 [=====] - 31s 309ms/step - loss: 0.3719 - accuracy: 0.9195 - val_loss: 0.1717 - val_accuracy: 0.9780
Epoch 6/10
100/100 [=====] - 30s 299ms/step - loss: 0.3617 - accuracy: 0.9145 - val_loss: 3.9468e-06 - val_accuracy: 0.9760
Epoch 7/10
100/100 [=====] - 31s 307ms/step - loss: 0.3225 - accuracy: 0.9245 - val_loss: 0.0853 - val_accuracy: 0.9710
Epoch 8/10
100/100 [=====] - 30s 299ms/step - loss: 0.3143 - accuracy: 0.9265 - val_loss: 1.0649e-05 - val_accuracy: 0.9780
Epoch 9/10
100/100 [=====] - 30s 297ms/step - loss: 0.3283 - accuracy: 0.9225 - val_loss: 0.3379 - val_accuracy: 0.9570
Epoch 10/10
100/100 [=====] - 30s 298ms/step - loss: 0.3275 - accuracy: 0.9220 - val_loss: 0.4117 - val_accuracy: 0.9790
```

Comparing The Four Models

| Name of Model | Number of Layers | Accuracy (%) | Validation Accuracy (%) | Loss (%) | Validation Loss (%) |
|---------------|------------------|--------------|-------------------------|----------|---------------------|
| Alex Net | 8 | 49.93 | 49.69 | 71.29 | 69.49 |
| VGG-16 | 16 | 84.75 | 90.45 | 34.18 | 23.2 |
| Efficient Net | 17 | 92.9 | 97.9 | 32.75 | 41.17 |
| Google Net | 22 | 91.45 | 96.2 | 31.36 | 13.2 |

Plan For Remaining Work

- Working on the MIAS dataset with the basic CNN and then proceeding with other models also.
- Trying to increase the accuracy and lower the losses for the MIAS dataset.



References

- [1] Krizhevsky Alex, Sutskever, I., H. Geoffrey E., "Alex Net," Advance Neural Inf Process. Syst, vol. 25, pp.1-9, 2012.
- [2] J. Fieres, J. Schemmel, and K. Meier, "Training convolutional networks of threshold neurons suited for low-power hardware implementation," In Neural Networks IJCNN'06. International Joint Conference on IEEE, pp. 21-28, 2012.
- [2] M. Lin, Q. Chen, and S. Yan, "Network in network," vol.1312, pp.4400, 2013.
- [3] J. G. Carbonell, R. S. Michalski, and T. M. Mitchell, "An overview of machine learning. In Machine learning." Springer Berlin Heidelberg , vol.3, pp.23, 2013.
- [4] C. Nebauer, "Evaluation of convolutional neural networks for visual recognition," IEEE Transactions on Neural Networks, vol.4, pp. 685-696, 2014.
- [5] Christian Szegedy, Wei Liu, Yangqing Jia, Pierre Sermanet, Scott Reed ,Dragomir Anguelov, Dumitru Erhan, Vincent Vanhoucke, Andrew Rabinovich, Google Inc. University of North Carolina, Chapel Hill, University of Michigan, Ann Arbor Magic Leap Inc. , "Going Deeper with Convolutions", 2014.

References

- [6] Karen Simonyan & Andrew Zisserman Visual Geometry Group, Department of Engineering Science, University of Oxford, "Very Deep Convolutional Networks For Large-Scale Image Recognition", 2014.
- [7] S. C. Dharmadhikari, M. Ingle, and P. Kulkarni, "Empirical studies on machine learning based text classification algorithms." Advanced Computing 2, vol.6, pp.161, 2015.
- [7] S. Albawi, T. A. Mohammed and S. Al-Zawi, "Understanding of a convolutional neural network," 2017 International Conference on Engineering and Technology (ICET), Antalya, pp. 1-6, 2017.
- [10] Tan, Mingxing, and Quoc V. Le. "Efficientnet: Rethinking model scaling for convolutional neural networks.", vol.1905, pp.11946, 2019.

